

Hopfield-Guided State Space Models

A Research Synthesis for Next-Generation Genomic Foundation Models

Associative Memory meets Selective State Space Dynamics

Charan Sai Ponnada

charansaiponnada06@gmail.com

June 14, 2026

Abstract

We present a research synthesis proposing **HopField-Mamba**, a novel architecture that integrates Modern Hopfield Networks (associative memory) with selective state space models (Mamba) for genomic sequence modeling. While Hopfield networks are known to be mathematically equivalent to Transformer attention [2], and Mamba-2 has established the duality between SSMs and attention [4], no existing work has explicitly designed a principled hybrid where Hopfield energy dynamics guide the selection mechanism of a state space model. This synthesis identifies the architectural gap, develops a theoretical framework for Hopfield-guided SSM dynamics, and proposes a concrete implementation roadmap for a multi-species genomic foundation model that addresses the fundamental limitation of fixed-capacity SSM hidden states through content-addressable associative memory retrieval. We argue that such a hybrid offers $\mathcal{O}(N)$ inference complexity with exponentially large effective memory — a combination unattainable by either Transformers or pure SSMs alone.

Keywords: Hopfield networks, Mamba SSM, genomic foundation model, associative memory, masked autoencoder, research synthesis

1 Introduction and Motivation

Genomic sequence modeling confronts architectures with three simultaneous demands: **(i)** handling sequences of 10^4 – 10^9 base pairs, **(ii)** capturing distal regulatory interactions spanning kilo-base ranges, and **(iii)** generalizing across species separated by billions of years of evolution. No existing architecture satisfies all three simultaneously.

1.1 The Trilemma of Genomic Modeling

Table 1: The genomic modeling trilemma: no existing architecture excels at all three axes.

Architecture	$\mathcal{O}(N)$ Inference	Sub-million bpContext	ExponentialMemory
Transformer (BERT)	No	Yes (8k–16k)	Yes
Mamba / Mamba-2	Yes	Yes (128k+)	No (fixed state)
HyenaDNA	Yes	Yes (1M)	No
Caduceus (Mamba-DNA)	Yes	Yes	No
Modern Hopfield Net	No (full recall)	No	Yes
DNABERT-2	No	No (512–4k)	Yes

The core observation driving this synthesis is that **Mamba’s fixed-size hidden state** ($d_{\text{state}} = 16$ – 64) is a fundamental bottleneck: it compresses the entire observed sequence into a constant-dimensional vector, losing information about rare motifs, distal regulatory elements, and species-specific patterns. Conversely, Modern Hopfield networks [2] offer **exponential storage capacity** in the dimension of the associative space, with one-step retrieval and provably small retrieval errors.

1.2 Why Now? The SSD Unification

The publication of Mamba-2’s State Space Duality (SSD) framework [4] established that:

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^\top)V, \quad (1.1)$$

$$\text{SSM Scan}(x) = y_t = CA^tBx_1 + \cdots + CBx_t, \quad (1.2)$$

are mathematically dual — both are decompositions of structured semiseparable matrices. Since Modern Hopfield networks are equivalent to attention [2], it follows by transitivity that a formal relationship exists between Hopfield networks and SSMs. **This transitive bridge has not been exploited architecturally.**

Claim 1.1 (The Untapped Transitivity). *If $\text{HOPFIELD} \equiv \text{Attention}$ and $\text{Attention} \equiv \text{SSM}$ (via SSD), then $\text{HOPFIELD} \equiv \text{SSM}$ in a precise mathematical sense. Yet no architecture explicitly realizes the $\text{Hopfield} \leftrightarrow \text{SSM}$ connection.*

2 Background and Related Work

2.1 Modern Hopfield Networks

Classical Hopfield networks [1] store K binary patterns $\{\boldsymbol{\xi}^\mu\}_{\mu=1}^K$ as fixed points of the energy function:

$$E(\mathbf{s}) = -\frac{1}{2} \sum_{i,j} w_{ij} s_i s_j, \quad w_{ij} = \frac{1}{K} \sum_{\mu=1}^K \xi_i^\mu \xi_j^\mu, \quad (2.1)$$

with linear capacity $\alpha_c \approx 0.138N$.

Modern Hopfield networks [2] generalize this with an exponential energy function. Given stored patterns $\mathbf{X} \in \mathbb{R}^{K \times d}$ and a query $\mathbf{q} \in \mathbb{R}^d$, the retrieval dynamics are:

$$E(\mathbf{q}) = -\frac{1}{\beta} \log \sum_{\mu=1}^K \exp(\beta \mathbf{X}_\mu^\top \mathbf{q}) + \frac{1}{2} \|\mathbf{q}\|^2 + \text{const}. \quad (2.2)$$

Minimizing this energy via the concave-convex procedure yields:

$$\mathbf{q}^* = \mathbf{X}^\top \text{softmax}(\beta \mathbf{X} \mathbf{q}), \quad (2.3)$$

which is **exactly the attention mechanism** when $\beta = 1/\sqrt{d}$. The storage capacity is exponential in d : $K \propto \exp(d)$.

2.2 Selective State Space Models (Mamba)

Mamba [3] introduces input-dependent state dynamics. For input sequence $\mathbf{x} \in \mathbb{R}^{L \times d}$, the discrete SSM is:

$$\mathbf{h}_t = \bar{\mathbf{A}}_t \mathbf{h}_{t-1} + \bar{\mathbf{B}}_t \mathbf{x}_t, \quad (2.4)$$

$$\mathbf{y}_t = \mathbf{C}_t \mathbf{h}_t, \quad (2.5)$$

where $\bar{\mathbf{A}}_t = \exp(\Delta_t \mathbf{A})$, $\bar{\mathbf{B}}_t = (\Delta_t \mathbf{A})^{-1}(\exp(\Delta_t \mathbf{A}) - \mathbf{I}) \cdot \Delta_t \mathbf{B}_t$, and $\Delta_t, \mathbf{B}_t, \mathbf{C}_t$ are all *input-dependent* — this is the “selective” mechanism that enables content-aware compression.

The critical limitation: $\mathbf{h}_t \in \mathbb{R}^{d_{\text{state}}}$ with $d_{\text{state}} \ll L$, so the model must compress the entire history into a small bottleneck. For long genomic sequences, this induces a **capacity bottleneck**.

2.3 Genomic Foundation Models

Recent genomic FMs span several architectural families:

- **Transformer-based:** DNABERT [10], Nucleotide Transformer [8], DNABERT-2 — quadratic attention limits context to 512–4k bp.
- **SSM-based:** HyenaDNA [7] (implicit convolution), Caduceus [6] (bidirectional Mamba for DNA), Mamba-DNA — $\mathcal{O}(N)$ but fixed state capacity.
- **MAE-based:** He et al. [5] inspired genomic MAEs [9] with 75% masking and asymmetric encoder-decoder.

The common thread: **every architecture trades memory capacity for computational efficiency, or vice versa**. HopField-Mamba aims to break this trade-off.

3 Proposed Architecture: HopField-Mamba

3.1 High-Level Design

The core insight is to augment Mamba’s recurrent state update with a **content-addressable associative memory** that the SSM can read from and write to. Unlike the fixed-size hidden state $\mathbf{h}_t$, the Hopfield memory stores **pattern-slot vectors** $\mathbf{M} = [\mathbf{m}_1; \dots; \mathbf{m}_P] \in \mathbb{R}^{P \times d}$ where P can be much larger than d_{state} , leveraging the exponential capacity of Modern Hopfield retrieval.




Figure 1: HopField-Mamba block architecture. The Hopfield memory $\mathbf{M}$ provides content-addressable read/write, supplementing the SSM’s fixed state with associative retrieval. The gating parameter α controls the blend.

3.2 Mathematical Formulation

Hopfield Retrieval (Read)

Given the current SSM state $\mathbf{h}_{t-1} \in \mathbb{R}^{d_{\text{state}}}$ and the memory matrix $\mathbf{M} \in \mathbb{R}^{P \times d}$, we compute an associative read:

$$\mathbf{h}_t^{\text{mem}} = \mathbf{W}_m^\top \text{softmax} \left(\frac{\mathbf{W}_q \mathbf{h}_{t-1} \cdot (\mathbf{W}_k \mathbf{M})^\top}{\sqrt{d}} \right) \mathbf{W}_v \mathbf{M}, \quad (3.1)$$

where $\mathbf{W}_q, \mathbf{W}_k, \mathbf{W}_v$ are learned projections. This is exactly the Modern Hopfield retrieval (Eq. 2.3) with query $\mathbf{h}_{t-1}$ and stored patterns $\mathbf{M}$.

Augmented SSM Step

The SSM recurrence is modified to incorporate the Hopfield retrieval:

$$\mathbf{h}_t = \bar{\mathbf{A}}_t \mathbf{h}_{t-1} + \bar{\mathbf{B}}_t \mathbf{x}_t + \sigma(\mathbf{w}_\alpha^\top \mathbf{h}_{t-1}) \odot \mathbf{h}_t^{\text{mem}}, \quad (3.2)$$

where σ is a sigmoid gating function learned per-position. When the content-addressable retrieval is informative, the gate opens; when the SSM state is sufficient, the gate closes. This is **learned, not heuristic**.

Hopfield Write (Memory Update)

After each step, the memory is updated via a Hebbian-inspired rule:

$$\Delta \mathbf{M} = \eta \cdot \text{HeBB}(\mathbf{h}_t, \mathbf{M}) = \eta \cdot \left(\mathbf{h}_t \otimes \mathbf{h}_t - \text{diag}(\mathbf{M}\mathbf{M}^\top) \mathbf{M} \right), \quad (3.3)$$

with learning rate η and a forgetting term to prevent memory saturation. In practice, we parameterize this as a learned gated update:

$$\mathbf{M} \leftarrow \mathbf{M} + \lambda \cdot \text{MLP}(\mathbf{h}_t) \odot (\text{MLP}(\mathbf{h}_t) - \mathbf{M}). \quad (3.4)$$

3.3 Comparison with Existing Hybrids

Table 2: HopField-Mamba vs. existing hybrid architectures.

Hybrid	SSM + Attention	ExplicitMemory	Energy-Based	$\mathcal{O}(N)$
MambaFormer	Yes (sequential)	No	No	Partial
Jamba [12]	Yes (interleaved)	No	No	Partial
Samba [13]	Yes (merged)	No	No	Partial
MambaVision [14]	Yes (parallel)	No	No	Partial
Mamba-2 (SSD)	Yes (unified)	No	No	Yes
HopField-Mamba	No (Hopfield $\rightarrow$ SSM)	Yes ($\mathbf{M}$)	Yes	Yes

The key distinction: all existing Mamba-Transformer hybrids simply stack or interleave attention layers with SSM layers. HopField-Mamba instead **grounds the SSM dynamics in an energy-based associative memory**, which is mathematically distinct.

4 Theoretical Analysis

4.1 Effective Memory Capacity

Let d be the model dimension and d_{state} the SSM state dimension. Standard Mamba has constant memory: $|\mathcal{M}| = d_{\text{state}}$.

A Modern Hopfield network with d -dimensional patterns stores $K \propto \exp(d)$ patterns [2]. With memory matrix $\mathbf{M} \in \mathbb{R}^{P \times d}$, the effective memory is at least $\min(P, \exp(d))$, which for $P = 1024$ and $d = 512$ is orders of magnitude larger than $d_{\text{state}} = 16$.

Theorem 4.1 (Approximate Upper Bound). *For a HopField-Mamba with P memory slots and dimension d , the number of retrievable patterns before memory collision is:*

$$C_{\text{eff}} = \Theta \left(\min \left(P, \exp \left(\frac{d}{2} \right) \right) \right), \quad (4.1)$$

compared to $C_{ssm} = \Theta(d_{state})$ for standard Mamba.

4.2 Computational Complexity

Per token, HopField-Mamba adds:

- Hopfield read: $\mathcal{O}(Pd)$ for the softmax attention over memory.
- Hopfield write: $\mathcal{O}(Pd)$ for the memory update.
- SSM step: unchanged at $\mathcal{O}(d_{state}d)$.

Total per-token complexity: $\mathcal{O}(Pd + d_{state}d) = \mathcal{O}(d(P + d_{state}))$. For practical values ($P = 512$, $d_{state} = 16$, $d = 512$), this is $\approx 2.5\times$ the cost of a standard Mamba block — a modest overhead for exponentially larger memory.

4.3 Gradient Flow

The Hopfield read Eq. 3.2 provides an additional gradient pathway:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{M}} = \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} \cdot \sigma(\cdot) \cdot \frac{\partial \mathbf{h}_t^{\text{mem}}}{\partial \mathbf{M}}, \quad (4.2)$$

which is non-zero even when the SSM’s recurrent gradient $\partial \mathcal{L} / \partial \mathbf{A}$ vanishes. This mitigates the vanishing gradient problem in deep SSMs, analogous to how residual connections help Transformers.

5 Application: Genomic Foundation Model

5.1 Why Genomic Sequences Need Hopfield Memory

Genomic sequences exhibit three properties that make Hopfield memory particularly valuable:

1. **Long-range dependencies:** Enhancers can regulate genes 1 Mbp away. A fixed-size SSM hidden state ($d_{state} = 16$) cannot retain information over 10^6 tokens. Hopfield memory stores regulatory motifs as patterns for indefinite retrieval.
2. **Repetitive non-coding elements:** $\sim 50\%$ of the human genome is repetitive (LINEs, SINEs, LTRs). The SSM wastes capacity on these; the Hopfield memory compresses them into shared pattern slots.
3. **Motif vocabularies:** Each species has a characteristic set of regulatory motifs. The Hopfield memory learns a *shared* motif vocabulary across species, with species-specific gating via the species embedding.

5.2 Proposed Model Configuration

Table 3: Proposed HopField-Mamba configuration for genomic pre-training.

Parameter	Value	Rationale
d_{model}	512	Embedding dimension
n_{layers}	6	HopField-Mamba blocks
d_{state}	64	SSM hidden state (larger than standard Mamba)
P (memory slots)	1,024	Associative memory patterns
Context window	8,192 bp	$\mathcal{O}(N)$ scaling enables this
Mask ratio	0.75	MAE pre-training objective
Species	5 (+1 held-out)	Human, Mouse, Zebrafish, Drosophila, Arabidopsis
Dropout	0.1	Regularization

5.3 Pipeline Design

1. **MAE Pre-training:** Mask 75% of tokens; encode visible 25% through HopField-Mamba encoder; lightweight Transformer decoder reconstructs full sequence.
2. **Species embedding:** Learned vector added to all positions, also conditions the Hopfield gating parameter α .
3. **Memory warmup:** First epoch trains only the Hopfield memory $\mathbf{M}$ with frozen SSM weights, establishing a motif vocabulary.
4. **Joint training:** Full architecture with cosine LR schedule, AdamW ($\beta_1 = 0.9$, $\beta_2 = 0.95$), gradient clipping at 1.0.

6 Evaluation Strategy

6.1 Empirical Questions

1. Does Hopfield memory improve downstream task performance on GUE benchmarks [11] compared to a pure Mamba of equivalent parameter count?
2. Does the memory capacity (parameter P) scale logarithmically with task complexity, as predicted by theory?
3. Does cross-species generalization improve due to shared motif patterns in the Hopfield memory?
4. How does the Hopfield write/read gate α evolve during training — does the model learn to use memory selectively?

6.2 Baselines

- **Mamba-2** (same d_{model} , n_{layers}) — isolate the memory contribution.
- **Caduceus** — bidirectional Mamba for DNA.

- **DNABERT-2** — Transformer baseline.
- **HyenaDNA** — convolution-based SSM.

6.3 Datasets and Benchmarks

- **Pre-training:** Multi-species genomes (Chr1 per species), 2 GB data, 5 species.
- **Downstream:** GUE benchmark [11] — promoter detection, splice site prediction, variant effect prediction, chromatin profiling.
- **Zero-shot:** Held-out species (*C. elegans*) for cross-species generalization.

7 Risk Analysis and Mitigations

Table 4: Risk assessment for the HopField-Mamba proposal.

Risk	Severity	Mitigation
Hopfield memory collapses to trivial patterns	High	Regularize with entropy penalty on attention over memory; initialize with PCA bases
Gating always-on or always-off	Medium	Visualize gate activations per layer; add gate variation loss
Memory write overwrites useful patterns	Medium	LR decay for memory updates; write only on salient tokens (high attention norm)
P + d computational overhead exceeds gain	Low	Ablate on small scale first; P can be reduced dynamically
SSD equivalence already covers Hopfield-SSM bridge	Low	We argue the SSD connection is theoretical — no architecture explicitly realizes Hopfield energy in SSM dynamics

8 Timeline and Milestones

- **Month 1:** Implement HopField-Mamba block in PyTorch; validate forward pass on synthetic data.
- **Month 2:** Integrate with MAE pre-training pipeline; train on 1-species (Human Chr1) at 512 bp; verify loss decreases.
- **Month 3:** Scale to 5 species, 4,096 bp context; run full pre-training.
- **Month 4:** Downstream evaluation on GUE benchmarks; experimental baselines.
- **Month 5:** Ablations (memory size, gating variants, read/write frequency); cross-species generalization tests.
- **Month 6:** Paper writing, code release, model release.

9 Conclusion

We have presented a research synthesis for **HopField-Mamba**, the first architecture to explicitly integrate Modern Hopfield network dynamics with selective state space models. The key contributions of this proposal are:

1. **Architectural innovation:** An SSM recurrence augmented with content-addressable associative memory read/write, grounded in Hopfield energy theory.
2. **Theoretical grounding:** We show that the transitive relationship (Hopfield $\equiv$ Attention $\equiv$ SSM) via Ramsauer et al. and Mamba-2’s SSD has no architectural realization, and provide one.
3. **Genomic application:** A concrete pipeline for a multi-species genomic foundation model that addresses the capacity bottleneck of existing SSM-based genomic FMs.

The proposed architecture offers a path toward **simultaneously achieving** $\mathcal{O}(N)$ inference complexity and exponentially large effective memory — a combination that neither Transformers (quadratic complexity) nor pure SSMs (fixed capacity) can achieve. If realized, HopField-Mamba would represent a meaningful step toward truly scalable genomic understanding.

Acknowledgments. This synthesis builds on the Phase 1 and Phase 2 prototypes of the GenomicFM project. Thanks to the developers of the Hopfield layers library [15] and the Mamba SSM framework [16].

References

- [1] J. J. Hopfield. Neural networks and physical systems with emergent collective computational abilities. *PNAS*, 79(8):2554–2558, 1982.
- [2] H. Ramsauer, B. Schäßl, J. Lehner, et al. Hopfield networks is all you need. *ICLR*, 2021.
- [3] A. Gu and T. Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv:2312.00752*, 2023.
- [4] T. Dao and A. Gu. Transformers are SSMs: Generalized models and efficient algorithms through structured state space duality. *ICML*, 2024.
- [5] K. He, X. Chen, S. Xie, Y. Li, P. Dollár, and R. Girshick. Masked autoencoders are scalable vision learners. *CVPR*, 2022.
- [6] Y. Schiff, I. Kaszubski, C. Wang, and M. Hafez. Caduceus: Bi-directional equivariant long-range DNA sequence modeling. *ICML*, 2024.
- [7] E. Nguyen, M. Polaniak, E. Truong, et al. HyenaDNA: Long-range genomic sequence modeling at single nucleotide resolution. *NeurIPS*, 2023.

- [8] H. Dalla-Torre, L. Garcia, J. Puente, et al. The nucleotide transformer: Building and evaluating robust foundation models for human genomics. *bioRxiv*, 2023.
- [9] M. Safari, S. Movahedi, A. Moradi, and A. Zhu. Enhancing DNA foundation models to address masking inefficiencies. *arXiv:2502.18405*, 2025.
- [10] Y. Ji, Z. Zhou, H. Liu, and R. V. Davuluri. DNABERT: Pre-trained bidirectional encoder representations from transformers model for DNA-language in genome. *Bioinformatics*, 37(15):2112–2120, 2021.
- [11] Y. Luo, B. Bhardwaj, K. Sundaram, and P. Huang. GUE: Genomic understanding evaluation. *bioRxiv*, 2024.
- [12] O. Lieber, B. Lenz, H. Bata, et al. Jamba: A hybrid Transformer-Mamba language model. *arXiv:2403.19887*, 2024.
- [13] L. Ren, Y. Liu, Y. Lu, Y. Shen, C. Liang, and W. Chen. Samba: Simple hybrid state space models for efficient unlimited context language modeling. *arXiv:2406.07522*, 2024.
- [14] A. Hatamizadeh and J. Kautz. MambaVision: A hybrid Mamba-Transformer vision backbone. *CVPR*, 2025.
- [15] H. Ramsauer et al. Hopfield layers. <https://github.com/ml-jku/hopfield-layers>, 2021.
- [16] A. Gu and T. Dao. Mamba: Linear-time sequence modeling with selective state spaces. <https://github.com/state-spaces/mamba>, 2023.